

---

# Deep Learning Foundations

Winter term 2024/25

Sandipan Sikdar



# Supervised Learning

---

- We are given a set of training examples of the form  $\{(x^1, y^1), \dots, (x^n, y^n)\}$
- $x^i$  represents the feature vector for the  $i$ th training example
- $y^i$  represents the ground truth
- Goal: Find  $f_W: R^d \rightarrow R, f_W(x^i) \approx y^i$
- $W$ : model parameters

# Supervised Learning

---


- Goal: Find  $f_W: R^d \rightarrow R, f_W(x^i) \approx y^i$
- How do we find  $W$ ?

$$\min_W \frac{1}{n} \sum_{i=1}^n l(f_W(x^i), y^i)$$

- $l$  -> loss/cost function
- Empirical risk minimization

# Choice of model



- Non-linear functions



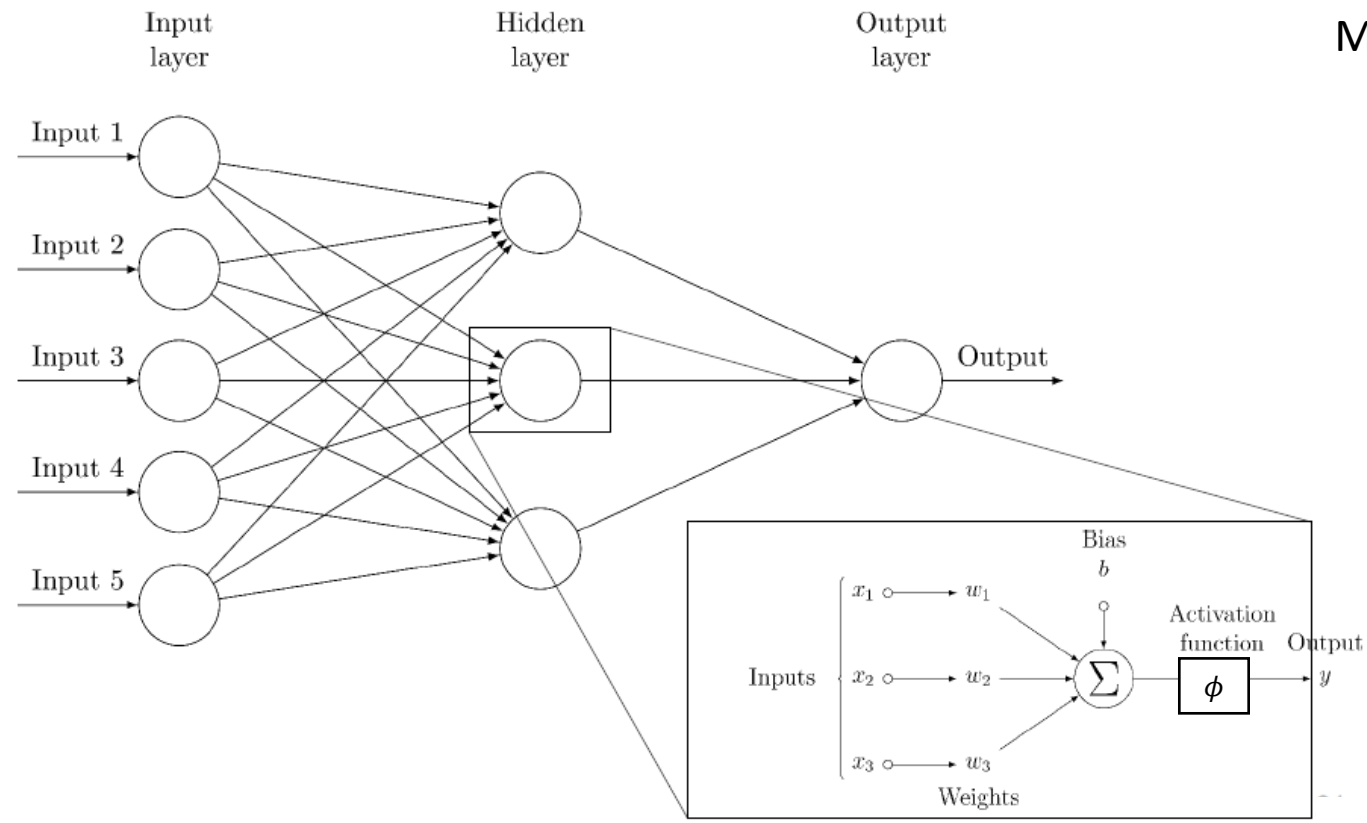
Simplest neural network

# Neural Network


- Cascaded layers



# Neural Network




# Neural Network – Forward Pass



$$h = \phi(W_1x + b_1)$$
$$y = \phi(W_2h + b_2)$$

# Neural Network – Forward Pass





# Chained Linear Classifiers

---




- Remember
  - A linear classifier provides class scores by calculating  $y = \phi(Wx + b)$
  - A neural network is a chain of linear classifiers with activation functions
- A neural network is some function like this:

$$y = \phi(W_3\phi(W_2\phi(W_1x + b_1) + b_2) + b_3)$$

- We can chain this as deep as we want

# Activation Functions

- Some common activation functions:



Some more:

- Heaviside
- Leaky ReLU
- SeLU
- ...

# Universal Approximation Theorem

---

- “A neural network can approximate any continuous function”

**Theorem 1.** *Let  $\sigma$  be any continuous discriminatory function. Then finite sums of the form*

$$G(x) = \sum_{j=1}^N \alpha_j \sigma(\overset{w_j}{\downarrow} y_j^T x + \theta_j) \quad (2)$$

*are dense in  $C(I_n)$ . In other words, given any  $f \in C(I_n)$  and  $\varepsilon > 0$ , there is a sum,  $G(x)$ , of the above form, for which*

$$|G(x) - f(x)| < \varepsilon \quad \text{for all } x \in I_n.$$

Cybenko G, “Approximation by Superpositions of a Sigmoidal Function”

# Training/Optimization

---

$$\min_W \frac{1}{n} \sum_{i=1}^n l(f_W(x^i), y^i)$$

- Use a search algorithm that starts with some initial guess and repeatedly changes  $W$  such that the loss gets smaller and smaller until we reach a point where the loss is minimized
- Gradient descent algorithm

# Training/Optimization

---

$$L(W) = \frac{1}{n} \sum_{i=1}^n l(f_W(x^i), y^i)$$


- Gradient descent algorithm

$$W_j = W_j - \alpha \frac{\partial}{\partial W_j} L(W)$$

Learning rate

# Gradient Descent

---



<http://blog.datumbox.com/wp-content/uploads/2013/10/gradient-descent.png>




# Gradient Descent

---

## Gradient Descent


### Gradient

= The vector of all partial derivatives of a function




### Descent

= Finding a way towards the minimum of the function



# Gradient

- Gradient provides the direction of steepest ascent



$$f(x_1, x_2) = x_1 x_2$$

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \end{bmatrix} = \begin{bmatrix} x_2 \\ x_1 \end{bmatrix}$$



# Backpropagation


---

- We need to compute derivative of complex functions
- ... but these are just chained simple functions
- Recursively apply chain rule

# Chain rule

---

- Example



$$z = f(y)$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$

$$y = f(x)$$

$$z = \sin(x^2)$$

Lets assume  $y = x^2$ , then  
we have

$$z = \sin(y)$$

$$\frac{\partial z}{\partial y} = \frac{\partial}{\partial y} \sin(y) = \cos(y) = \cos(x^2)$$

$$\frac{\partial y}{\partial x} = \frac{\partial}{\partial x} x^2 = 2x$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \cos(x^2) 2x$$


# Chain rule (multi-variable)

---

- Given a multivariable function  $f(x, y)$ , and two single variable functions  $x(t)$  and  $y(t)$ , here's what the multivariable chain rule says:

$$\underbrace{\frac{d}{dt} f(x(t), y(t))}_{\text{Derivative of composition function}} = \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$

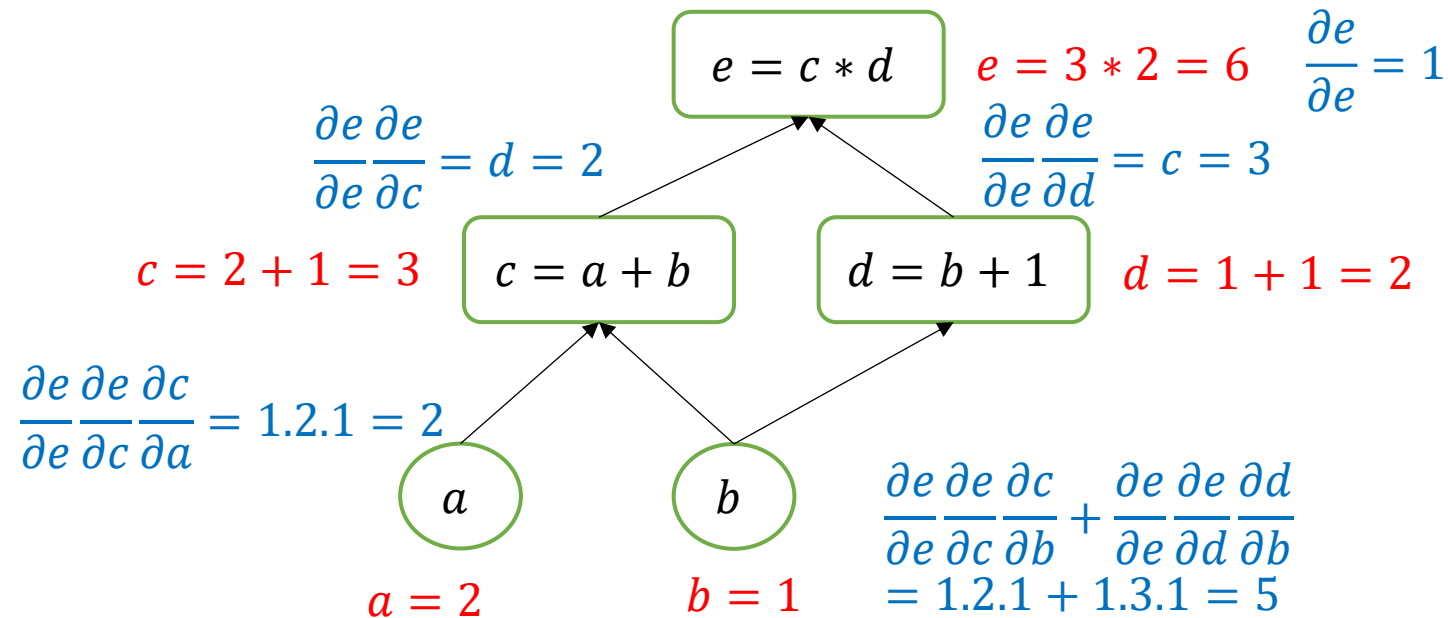
Derivative of composition function



# Computation graph

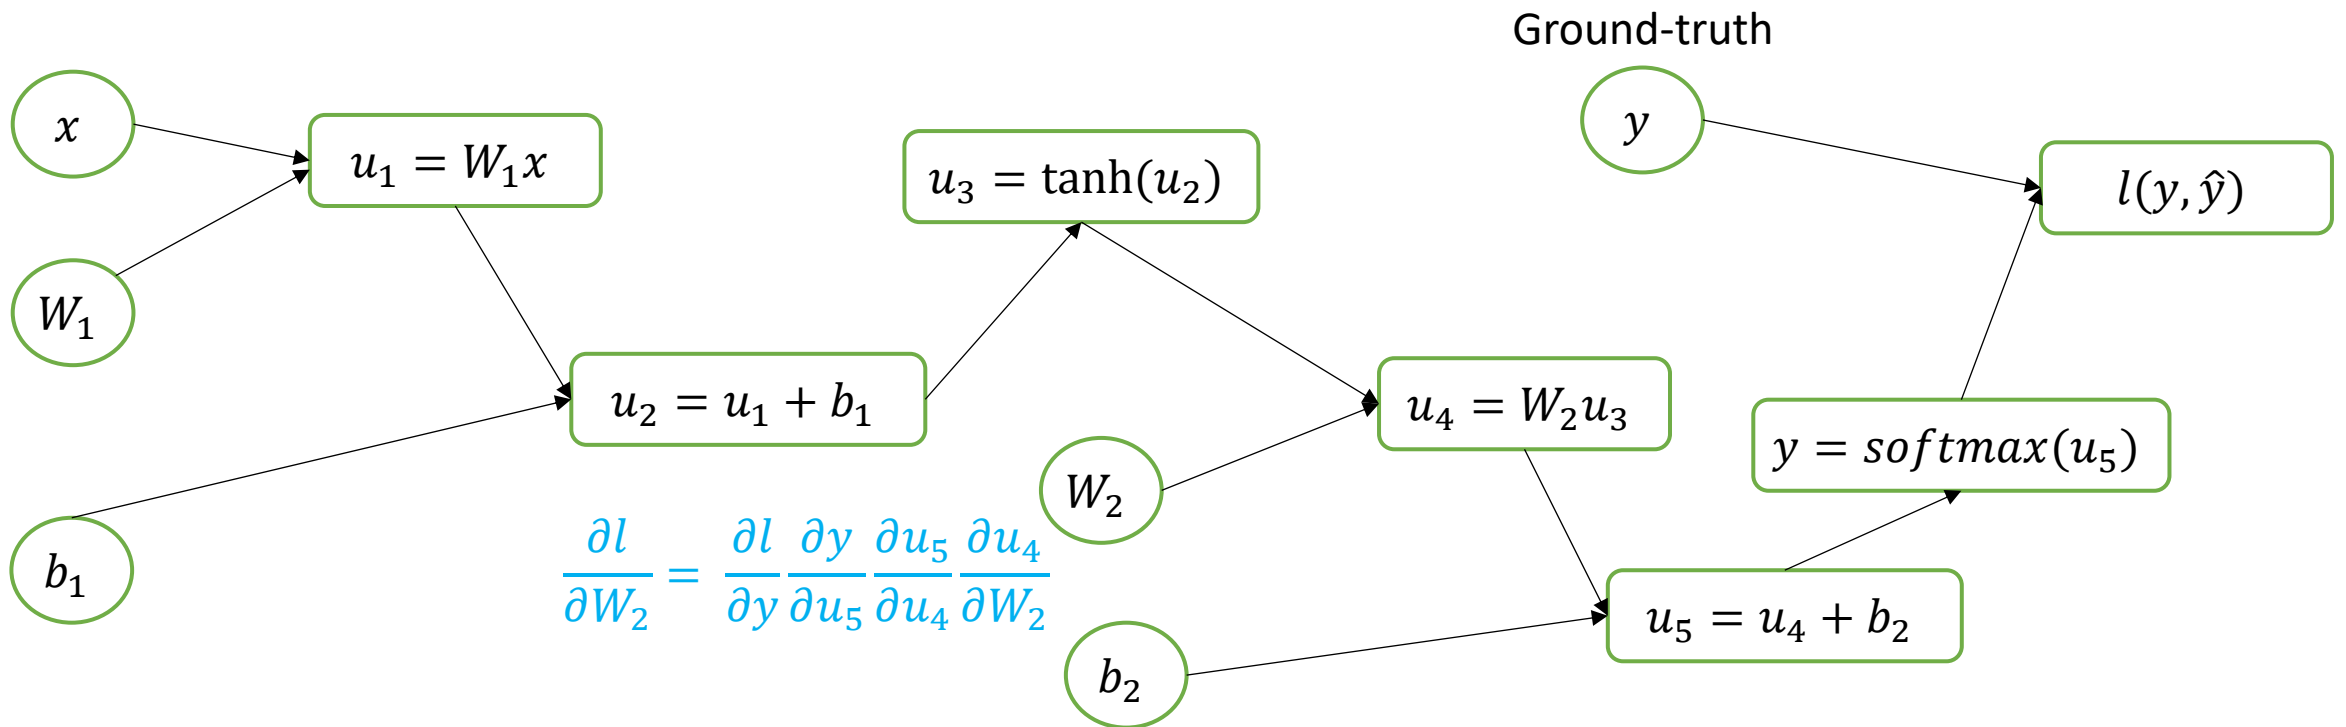
- The computations are organized as a graph

$$f(a, b) = (a + b) * (b + 1)$$



# Computation graph

$$\hat{y} = \text{softmax}(W_2 \tanh(w_1 x + b_1) + b_2)$$



# Gradient Descent


---

## Gradient Descent

**Gradient**




= The vector of all partial derivatives of a function



**Descent**

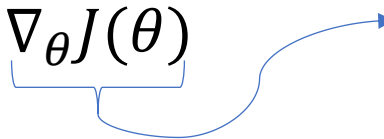
= Finding a way towards the minimum of the function



# Gradient Descent

---

- Size of steps = Learning rate ( $\eta$ )
- Parameters of the model =  $\theta$
- Loss computed on the entire dataset =  $J(\theta)$

$$\theta = \theta - \eta \nabla_{\theta} J(\theta) \quad \nabla_{\theta} J(\theta) = \frac{\partial}{\partial \theta} J(\theta)$$


- Compute over multiple epochs

```
for i in range(nb_epochs):  
    params_grad = evaluate_gradient(loss_function, data, params)  
    params = params - learning_rate * params_grad
```

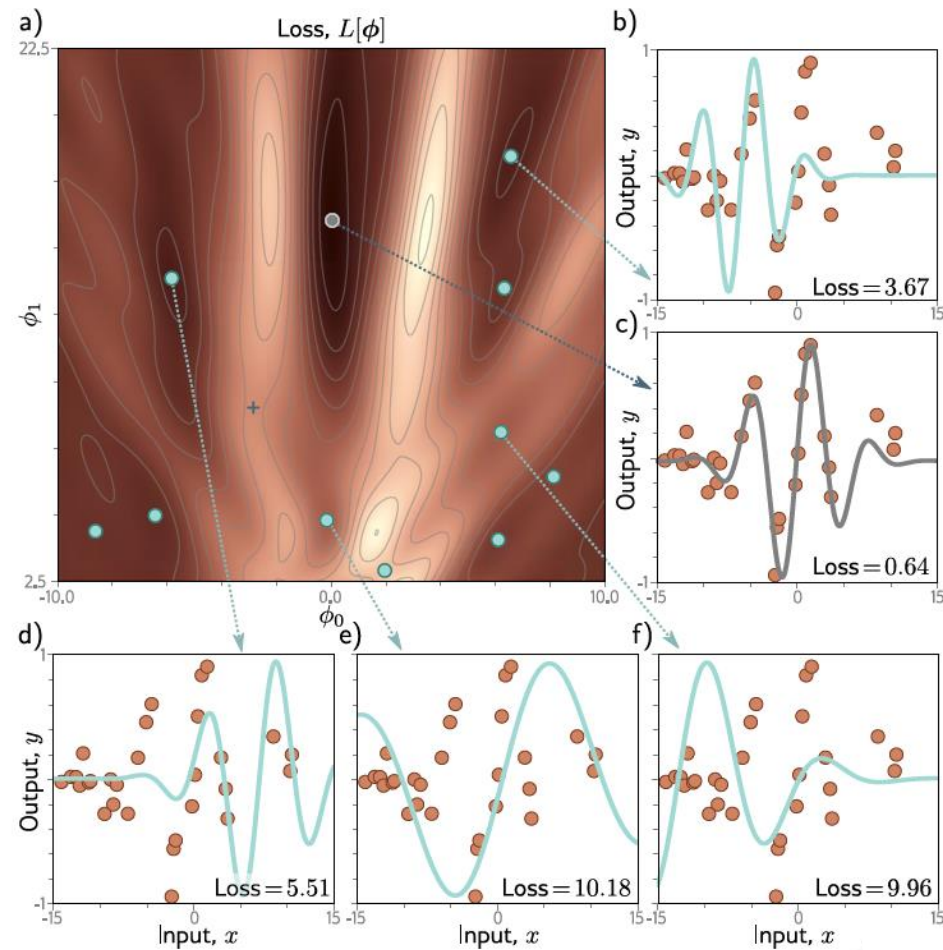
# Gradient Descent

---

- GD is guaranteed to reach minima if the optimization function is convex
- It does not matter how you initialize the parameters
- Unfortunately, loss functions for non-linear models are non-convex



# Local minima



# Stochastic Gradient Descent

---

- Size of steps = Learning rate ( $\eta$ )
- Parameters of the model =  $\theta$
- At each step calculate objective function over a randomly chosen training example  $(x^i, y^i)$
- Update for  $\theta$

$$\theta = \theta - \eta \nabla_{\theta} J(\theta; x^i, y^i)$$

```
for i in range(nb_epochs):  
    np.random.shuffle(data)  
    for example in data:  
        params_grad = evaluate_gradient(loss_function, example, params)  
        params = params - learning_rate * params_grad
```

# Minibatch Gradient Descent

---

- Size of steps = Learning rate ( $\eta$ )
- Parameters of the model =  $\theta$
- At each step calculate objective function over a minibatch of  $n$  training examples
- Update for  $\theta$

$$\theta = \theta - \eta \nabla_{\theta} J(\theta; x^{i:i+n}, y^{i:i+n})$$

```
for i in range(nb_epochs):  
    np.random.shuffle(data)  
    for batch in get_batches(data, batch_size=50):  
        params_grad = evaluate_gradient(loss_function, batch, params)  
        params = params - learning_rate * params_grad
```

# Momentum

---

- Momentum: Helps accelerate SGD
- It does so by adding a fraction  $\gamma$  of the update vector of the past time step to the current vector

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$$

$$\theta = \theta - v_t$$

# Nesterov Momentum

---

- Nesterov Momentum
  - Momentum tends to overshoot the minimum -> Prevent that!
  - Use a lookahead: Take the gradient at the position that would be reached by the next step with the current velocity
- We know we will use momentum term  $\gamma v_{t-1}$  to move the parameter  $\theta$
- $\theta - \gamma v_{t-1}$  gives us an approximation of where the parameter is going to be (lookahead)
- We calculate gradient with respect to our approximate future parameter instead of the current one

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$

$$\theta = \theta - v_t$$

# Adagrad

---

- Introduces a cache variable that modifies the effective learning rate per parameter  $\theta_i$

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} g_{t,i}$$

$\nabla_{\theta} J(\theta_t, i)$

Sum of squares of past gradients with respect to  $\theta$

Normalize learning rate based on history

# Adagrad

---

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} g_{t,i}$$

- Decreases learning rate for parameters that have large / frequent updates
- Increases learning rate for parameters that have small / infrequent updates
- Advantages:
  - NO (less) tuning of learning rates, (less sensitive to hyperparameter)
  - Faster convergence if scaling of the weights is unequal

# Adam

---

- Adagrad with momentum

$$\left. \begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \end{aligned} \right\} \text{Consider both first and second moments}$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

- $m_t$  and  $v_t$  are initialized to 0 and are hence biased towards 0 in the initial steps
- Correct for such biases -


$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$



# Learning rate

---


- SGD, SGD+momentum, Adagrad, Adam use learning rate as hyperparameter
- How do we choose a good learning rate?



# Learning rate decays over time

---

- Step: Reduce learning rate at a few fixed points
- E.g., multiply learning rate by 0.1 after 30, 60, 90 epochs




# Learning Rate Decay

---

- Cosine

$$\alpha_t = \frac{1}{2} \alpha_0 (1 + \cos(t\pi/T))$$

- $\alpha_0$ : Initial learning rate
- $\alpha_t$ : Learning rate at epoch  $t$
- $T$ : Total number of epochs




# Learning Rate Decay

---

- Linear

$$\alpha_t = \alpha_0(1 - t/T)$$

- $\alpha_0$ : Initial learning rate
- $\alpha_t$ : Learning rate at epoch  $t$
- $T$ : Total number of epochs




# Learning Rate Decay

---

- Inverted sqrt

$$\alpha_t = \alpha_0 / \sqrt{t}$$


- $\alpha_0$ : Initial learning rate
- $\alpha_t$ : Learning rate at epoch  $t$
- $T$ : Total number of epochs



# Learning Rate Decay: Linear Warmup

---

- High initial learning rates can make loss explode
- Linearly increasing the learning rate from 0 over first ~5000 iterations can prevent this



# Parameter Initialization

---

- How to initialize the parameters before we start training?

<https://www.deeplearning.ai/ai-notes/initialization/index.html>

# Parameter Initialization

---

- Initializing to 0 or to a constant
  - Hidden units will have identical influence on loss
  - Prevent different neurons from learning different things
- Initializing weights to very high values -> exploding gradient
- Initializing weights to very low values -> vanishing gradient



# Appropriate Initialization

---

- Rules of thumb
  - The mean of the activations should be zero.
  - The variance of the activations should stay the same across every layer.
- Ensuring zero-mean and maintaining the value of the variance of the input of every layer guarantees no exploding/vanishing signal

# Standard Normal

---

- Pick numbers from a standard normal distribution

```
W = 0.01*np.random.randn()
```

- More inputs lead to higher variance
- Can be fixed through normalization

$$w \sim N(0, 1/n)$$

- $n$  is the number of inputs to the neuron

# Xavier Initialization

---

- Introduced by Xavier Glorot and Yoshua Bengio
- Recommend to normalize the variance to

$$Var(w) = \frac{2}{n_{in} + n_{out}}$$

$w \sim N(0, Var(w))$

Number of outputs

Number of inputs

Works best with sigmoid or tanh activation functions

# Kaiming Initialization

---

- Introduced by Kaiming He et. Al

$$Var(w) = \frac{gain}{n_{in}}$$

- *gain* depends on the activation function (e.g., for ReLU, *gain*=2)


$$w \sim N(0, Var(w))$$

- Works best with ReLU

# Batch Normalization

---

- Explicitly make the outputs Gaussian



$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

# Batch Normalization

---

- Learnable scale and shift parameters:  $\gamma, \beta$

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}} \quad \Rightarrow \quad y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

# Batch Normalization

---

- During test time,  $\mu_j, \sigma_j$  are not computed, rather an estimate from training is used


$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$


# Layer Normalization

---

- Same as batch norm only that the estimates are computed along dimensions



$D$   
Batch Norm




$D$   
Layer Norm



# Regularization



- Generalization/Test error
  - Performance on previously unseen inputs



- Regularization is:
  - Any modification to a learning algorithm to reduce its generalization error but not its training error
  - Reduce generalization error at the expense of test error

# Regularization

---



# Regularization

---

$$L(W) = \frac{1}{n} \sum_{i=1}^n l(f_W(x^i), y^i) + \lambda R(W)$$

- Data loss: Model predictions should match training data
- Regularization: Prevent the model from doing too well on the training data

# Regularization

---

$$L(W) = \frac{1}{n} \sum_{i=1}^n l(f_W(x^i), y^i) + \lambda R(W)$$

- L2 regularization:

$$R(W) = \frac{1}{2} \sum |W|^2$$

- L1 regularization:

$$R(W) = \sum |W|$$

# Regularization

---

- Expressing preference over weights

$$x = [1, 1, 2, 1]$$

Which  $W$  will  $L_2$  prefer?

$$W_1 = [0, 0, 1, 0]$$

Which  $W$  will  $L_1$  prefer?

$$W_2 = [0.5, 0.5, 0.25, 0.5]$$

$$W_1^T x = W_2^T x = 2$$

# Regularization

---

- $L_1$  prefers weights which are sparse
- $L_2$  prefers weights which are more spread out
- Decide on which regularization to use depending on the task.
- $L_2$  is used more often

# Regularization

---

- Elastic net ( $L_1 + L_2$ )

$$R(W) = |W| + \beta |W|^2$$


- Weight decay: Adds a regularization term to the gradient of loss

$$\frac{\partial}{\partial W} L + \lambda W$$

# Regularization

---

- Dropout: In each forward pass, randomly set some neurons to Zero



- Force the model to not rely on only a certain set of features



# Regularization

---

- Dropout makes our output random

$$y = f_W(x, z) \text{ Random mask}$$

- Want to “average out” the randomness at test time


$$y = f(x) = E_z[f_W(x, z)] = \int p(z) f_W(x, z) dz$$

Difficult to compute...

# Regularization

---

- We would like to approximate the integral



$$\text{At test: } E[a] = W_1x + W_2y$$

During training:

$$\begin{aligned} E[a] &= \frac{1}{4}(W_1x + W_2y) + \frac{1}{4}(W_1x + 0y) + \frac{1}{4}(0x + 0y) + \frac{1}{4}(0x + W_2y) \\ &= \frac{1}{2}(W_1x + W_2y) \end{aligned}$$

At test time multiply with the dropout probability

# Regularization


---

- Inverted dropout
- Normalize by  $p$  during training
- Test time is unchanged

# Early Stopping

---

- Stop training the model when accuracy on the validation set decreases Or train for a long time, but always keep track of the model snapshot that worked best on val



# Summary

---

- Training neural networks
  - Optimization
  - Weight Initialization
  - Regularization

# References

---

- Regularization slides adopted from CS231n course at Stanford
- <https://www.ruder.io/optimizing-gradient-descent/> (Optimization)
- <https://colah.github.io/posts/2015-08-Backprop/>